

# Working with Data: Day 1

*May 3-4, 2018*

## Outline

- review data types in R
- data frame manipulation with dplyr
- indexing with base R

## Data types in R

- vectors
  - zero to many elements
  - all elements are the same type (and NA)
  - logical, integer, numeric, character, factor
  - 2D matrices and many-D arrays are still vectors, just with the dimensions encoded in the “dim” attribute
- lists
  - aka “generic vectors”
  - any element can be a vector of any type and length
- data.frames
  - just a list with one restriction: all elements have the same length
  - a general representation for a 2D table of data
  - each **list** element is a column in the table

## Data frame

- representation of a 2D table of data
- columns are variables (something you measure/observe)
- rows are observations (variable values) of a specific **case**

## Data manipulation

Typical tasks:

- view a subset of variables
- choose a subset of observations based on one or more variable values
- create new variables based on existing variables
- order observations by some variable value
- calculate a summary of a set of variable values

Base R can do all of those things, but it can be pretty low-level (focus on coding instead of analyzing data) and awkward, so a number of alternatives have been put forward over time.

Currently, the most popular such alternative solution for data manipulation is the package dplyr. It’s so good at what it does, and integrates so well with other popular tools like *ggplot2*, that it has rapidly become the de-facto standard and it is what we will focus on today.

Dplyr has a set of functions, or *verbs* in its terminology, that each deal with one of the above tasks:

- **select**: view only some variables

- **filter**: choose observations by their values
- **arrange**: order observations (rows)
- **mutate**: create new variables
- **summarise**: calculate a summary of many variable values

Each verb works similarly:

- input data frame in the first argument
- other arguments can refer to variables as if they were local objects
- output is another data frame

## Select

You can use the **select** function to focus on a subset of variables.

```
library(dplyr)
select(mtcars,mpg,wt)
```

```
##           mpg      wt
## Mazda RX4      21.0  2.620
## Mazda RX4 Wag  21.0  2.875
## Datsun 710     22.8  2.320
## Hornet 4 Drive  21.4  3.215
## Hornet Sportabout 18.7  3.440
## Valiant        18.1  3.460
## Duster 360     14.3  3.570
## Merc 240D      24.4  3.190
## Merc 230       22.8  3.150
## Merc 280       19.2  3.440
## Merc 280C      17.8  3.440
## Merc 450SE     16.4  4.070
## Merc 450SL     17.3  3.730
## Merc 450SLC    15.2  3.780
## Cadillac Fleetwood 10.4  5.250
## Lincoln Continental 10.4  5.424
## Chrysler Imperial 14.7  5.345
## Fiat 128       32.4  2.200
## Honda Civic    30.4  1.615
## Toyota Corolla 33.9  1.835
## Toyota Corona  21.5  2.465
## Dodge Challenger 15.5  3.520
## AMC Javelin    15.2  3.435
## Camaro Z28     13.3  3.840
## Pontiac Firebird 19.2  3.845
## Fiat X1-9      27.3  1.935
## Porsche 914-2  26.0  2.140
## Lotus Europa   30.4  1.513
## Ford Pantera L 15.8  3.170
## Ferrari Dino   19.7  2.770
## Maserati Bora   15.0  3.570
## Volvo 142E     21.4  2.780
```

There are many helpful functions that can be used with **select** to describe which variables to keep:

- **starts\_with**(x, ignore.case = TRUE): names starts with x
- **ends\_with**(x, ignore.case = TRUE): names ends in x

- `contains(x, ignore.case = TRUE)`: selects all variables whose name contains `x`
- `matches(x, ignore.case = TRUE)`: selects all variables whose name matches the regular expression `x`
- `num_range("x", 1:5, width = 2)`: selects all variables (numerically) from `x01` to `x05`.
- `one_of("x", "y", "z")`: selects variables provided in a character vector.
- `everything()`: selects all variables.

```
select(mtcars, starts_with("d"))
```

```
##           disp drat
## Mazda RX4      160.0 3.90
## Mazda RX4 Wag  160.0 3.90
## Datsun 710     108.0 3.85
## Hornet 4 Drive 258.0 3.08
## Hornet Sportabout 360.0 3.15
## Valiant        225.0 2.76
## Duster 360     360.0 3.21
## Merc 240D      146.7 3.69
## Merc 230       140.8 3.92
## Merc 280       167.6 3.92
## Merc 280C      167.6 3.92
## Merc 450SE     275.8 3.07
## Merc 450SL     275.8 3.07
## Merc 450SLC    275.8 3.07
## Cadillac Fleetwood 472.0 2.93
## Lincoln Continental 460.0 3.00
## Chrysler Imperial 440.0 3.23
## Fiat 128       78.7 4.08
## Honda Civic    75.7 4.93
## Toyota Corolla 71.1 4.22
## Toyota Corona 120.1 3.70
## Dodge Challenger 318.0 2.76
## AMC Javelin    304.0 3.15
## Camaro Z28     350.0 3.73
## Pontiac Firebird 400.0 3.08
## Fiat X1-9      79.0 4.08
## Porsche 914-2  120.3 4.43
## Lotus Europa   95.1 3.77
## Ford Pantera L 351.0 4.22
## Ferrari Dino   145.0 3.62
## Maserati Bora  301.0 3.54
## Volvo 142E     121.0 4.11
```

This trick is handy to reorder the variables so that the ones you're most interested in are at the front, without dropping any:

```
select(mtcars, cyl, everything())
```

```
##           cyl mpg  disp  hp drat    wt  qsec vs  am gear carb
## Mazda RX4      6 21.0 160.0 110 3.90 2.620 16.46 0  1    4    4
## Mazda RX4 Wag  6 21.0 160.0 110 3.90 2.875 17.02 0  1    4    4
## Datsun 710     4 22.8 108.0  93 3.85 2.320 18.61 1  1    4    1
## Hornet 4 Drive  6 21.4 258.0 110 3.08 3.215 19.44 1  0    3    1
## Hornet Sportabout 8 18.7 360.0 175 3.15 3.440 17.02 0  0    3    2
## Valiant        6 18.1 225.0 105 2.76 3.460 20.22 1  0    3    1
## Duster 360     8 14.3 360.0 245 3.21 3.570 15.84 0  0    3    4
## Merc 240D      4 24.4 146.7  62 3.69 3.190 20.00 1  0    4    2
```

## Merc 230	4	22.8	140.8	95	3.92	3.150	22.90	1	0	4	2
## Merc 280	6	19.2	167.6	123	3.92	3.440	18.30	1	0	4	4
## Merc 280C	6	17.8	167.6	123	3.92	3.440	18.90	1	0	4	4
## Merc 450SE	8	16.4	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	8	17.3	275.8	180	3.07	3.730	17.60	0	0	3	3
## Merc 450SLC	8	15.2	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	8	10.4	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	8	10.4	460.0	215	3.00	5.424	17.82	0	0	3	4
## Chrysler Imperial	8	14.7	440.0	230	3.23	5.345	17.42	0	0	3	4
## Fiat 128	4	32.4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Honda Civic	4	30.4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	4	33.9	71.1	65	4.22	1.835	19.90	1	1	4	1
## Toyota Corona	4	21.5	120.1	97	3.70	2.465	20.01	1	0	3	1
## Dodge Challenger	8	15.5	318.0	150	2.76	3.520	16.87	0	0	3	2
## AMC Javelin	8	15.2	304.0	150	3.15	3.435	17.30	0	0	3	2
## Camaro Z28	8	13.3	350.0	245	3.73	3.840	15.41	0	0	3	4
## Pontiac Firebird	8	19.2	400.0	175	3.08	3.845	17.05	0	0	3	2
## Fiat X1-9	4	27.3	79.0	66	4.08	1.935	18.90	1	1	4	1
## Porsche 914-2	4	26.0	120.3	91	4.43	2.140	16.70	0	1	5	2
## Lotus Europa	4	30.4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Ford Pantera L	8	15.8	351.0	264	4.22	3.170	14.50	0	1	5	4
## Ferrari Dino	6	19.7	145.0	175	3.62	2.770	15.50	0	1	5	6
## Maserati Bora	8	15.0	301.0	335	3.54	3.570	14.60	0	1	5	8
## Volvo 142E	4	21.4	121.0	109	4.11	2.780	18.60	1	1	4	2

Using a named argument will rename a variable:

```
select(mtcars, mpg, weight=wt)
```

##	mpg	weight
## Mazda RX4	21.0	2.620
## Mazda RX4 Wag	21.0	2.875
## Datsun 710	22.8	2.320
## Hornet 4 Drive	21.4	3.215
## Hornet Sportabout	18.7	3.440
## Valiant	18.1	3.460
## Duster 360	14.3	3.570
## Merc 240D	24.4	3.190
## Merc 230	22.8	3.150
## Merc 280	19.2	3.440
## Merc 280C	17.8	3.440
## Merc 450SE	16.4	4.070
## Merc 450SL	17.3	3.730
## Merc 450SLC	15.2	3.780
## Cadillac Fleetwood	10.4	5.250
## Lincoln Continental	10.4	5.424
## Chrysler Imperial	14.7	5.345
## Fiat 128	32.4	2.200
## Honda Civic	30.4	1.615
## Toyota Corolla	33.9	1.835
## Toyota Corona	21.5	2.465
## Dodge Challenger	15.5	3.520
## AMC Javelin	15.2	3.435
## Camaro Z28	13.3	3.840
## Pontiac Firebird	19.2	3.845

```
## Fiat X1-9          27.3  1.935
## Porsche 914-2     26.0  2.140
## Lotus Europa      30.4  1.513
## Ford Pantera L    15.8  3.170
## Ferrari Dino      19.7  2.770
## Maserati Bora     15.0  3.570
## Volvo 142E       21.4  2.780
```

(You can also use `rename()` to change variable names while keeping all columns as they were.)

## Filter

You can use `filter` to select specific rows based on a logical condition of a variable. To specify more than one condition, just give them as additional arguments. The conditions are joined together as a logical *and*:

```
filter(mtcars, cyl==8)
```

```
##      mpg  cyl  disp  hp drat    wt  qsec vs am gear carb
## 1  18.7    8 360.0 175 3.15 3.440 17.02 0  0    3    2
## 2  14.3    8 360.0 245 3.21 3.570 15.84 0  0    3    4
## 3  16.4    8 275.8 180 3.07 4.070 17.40 0  0    3    3
## 4  17.3    8 275.8 180 3.07 3.730 17.60 0  0    3    3
## 5  15.2    8 275.8 180 3.07 3.780 18.00 0  0    3    3
## 6  10.4    8 472.0 205 2.93 5.250 17.98 0  0    3    4
## 7  10.4    8 460.0 215 3.00 5.424 17.82 0  0    3    4
## 8  14.7    8 440.0 230 3.23 5.345 17.42 0  0    3    4
## 9  15.5    8 318.0 150 2.76 3.520 16.87 0  0    3    2
## 10 15.2    8 304.0 150 3.15 3.435 17.30 0  0    3    2
## 11 13.3    8 350.0 245 3.73 3.840 15.41 0  0    3    4
## 12 19.2    8 400.0 175 3.08 3.845 17.05 0  0    3    2
## 13 15.8    8 351.0 264 4.22 3.170 14.50 0  1    5    4
## 14 15.0    8 301.0 335 3.54 3.570 14.60 0  1    5    8
```

```
filter(mtcars, cyl==8, carb==3)
```

```
##      mpg  cyl  disp  hp drat    wt  qsec vs am gear carb
## 1  16.4    8 275.8 180 3.07 4.07 17.4  0  0    3    3
## 2  17.3    8 275.8 180 3.07 3.73 17.6  0  0    3    3
## 3  15.2    8 275.8 180 3.07 3.78 18.0  0  0    3    3
```

To use the logical *or* to join conditions, you must use the `|` operator explicitly:

```
filter(mtcars, cyl==4 | carb==8)
```

```
##      mpg  cyl  disp  hp drat    wt  qsec vs am gear carb
## 1  22.8    4 108.0  93 3.85 2.320 18.61 1  1    4    1
## 2  24.4    4 146.7  62 3.69 3.190 20.00 1  0    4    2
## 3  22.8    4 140.8  95 3.92 3.150 22.90 1  0    4    2
## 4  32.4    4  78.7  66 4.08 2.200 19.47 1  1    4    1
## 5  30.4    4  75.7  52 4.93 1.615 18.52 1  1    4    2
## 6  33.9    4  71.1  65 4.22 1.835 19.90 1  1    4    1
## 7  21.5    4 120.1  97 3.70 2.465 20.01 1  0    3    1
## 8  27.3    4  79.0  66 4.08 1.935 18.90 1  1    4    1
## 9  26.0    4 120.3  91 4.43 2.140 16.70 0  1    5    2
## 10 30.4    4  95.1 113 3.77 1.513 16.90 1  1    5    2
## 11 15.0    8 301.0 335 3.54 3.570 14.60 0  1    5    8
```

```
## 12 21.4 4 121.0 109 4.11 2.780 18.60 1 1 4 2
```

```
filter(mtcars, cyl==8 & carb==8)
```

```
## mpg cyl disp hp drat wt qsec vs am gear carb
```

```
## 1 15 8 301 335 3.54 3.57 14.6 0 1 5 8
```

```
mtcars[mtcars$cyl==4 | mtcars$carb==8, ]
```

```
##      mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## Datsun 710  22.8  4 108.0  93 3.85 2.320 18.61 1 1  4  1
## Merc 240D  24.4  4 146.7  62 3.69 3.190 20.00 1 0  4  2
## Merc 230   22.8  4 140.8  95 3.92 3.150 22.90 1 0  4  2
## Fiat 128   32.4  4  78.7  66 4.08 2.200 19.47 1 1  4  1
## Honda Civic 30.4  4  75.7  52 4.93 1.615 18.52 1 1  4  2
## Toyota Corolla 33.9  4  71.1  65 4.22 1.835 19.90 1 1  4  1
## Toyota Corona 21.5  4 120.1  97 3.70 2.465 20.01 1 0  3  1
## Fiat X1-9   27.3  4  79.0  66 4.08 1.935 18.90 1 1  4  1
## Porsche 914-2 26.0  4 120.3  91 4.43 2.140 16.70 0 1  5  2
## Lotus Europa 30.4  4  95.1 113 3.77 1.513 16.90 1 1  5  2
## Maserati Bora 15.0  8 301.0 335 3.54 3.570 14.60 0 1  5  8
## Volvo 142E  21.4  4 121.0 109 4.11 2.780 18.60 1 1  4  2
```

If you need to select several conditions on the same variable you can use `%in%`:

```
filter(mtcars, carb==3 | carb==6 | carb==8)
```

```
##      mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## 1 16.4  8 275.8 180 3.07 4.07 17.4 0 0  3  3
## 2 17.3  8 275.8 180 3.07 3.73 17.6 0 0  3  3
## 3 15.2  8 275.8 180 3.07 3.78 18.0 0 0  3  3
## 4 19.7  6 145.0 175 3.62 2.77 15.5 0 1  5  6
## 5 15.0  8 301.0 335 3.54 3.57 14.6 0 1  5  8
```

```
filter(mtcars, carb %in% c(3,6,8))
```

```
##      mpg cyl  disp  hp drat   wt  qsec vs am gear carb
## 1 16.4  8 275.8 180 3.07 4.07 17.4 0 0  3  3
## 2 17.3  8 275.8 180 3.07 3.73 17.6 0 0  3  3
## 3 15.2  8 275.8 180 3.07 3.78 18.0 0 0  3  3
## 4 19.7  6 145.0 175 3.62 2.77 15.5 0 1  5  6
## 5 15.0  8 301.0 335 3.54 3.57 14.6 0 1  5  8
```

To use numeric indices the *dplyr* function is `slice`.

```
slice(mtcars, c(1,3,5))
```

```
## # A tibble: 3 x 11
```

```
##      mpg   cyl  disp    hp  drat    wt  qsec    vs  am  gear  carb
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  21.0     6.  160.   110.   3.90  2.62  16.5     0.    1.    4.    4.
## 2  22.8     4.  108.    93.   3.85  2.32  18.6     1.    1.    4.    1.
## 3  18.7     8.  360.   175.   3.15  3.44  17.0     0.    0.    3.    2.
```

Note: `slice` and `filter` do not carry the row names with the subset of rows.

If you wish to include the row names you need to add them to the data frame as a variable

```
filter(add_rownames(mtcars), cyl==8, carb==3)
```

```
## Warning: Deprecated, use tibble::rownames_to_column() instead.
## # A tibble: 3 x 12
##   rowname      mpg   cyl  disp    hp  drat    wt   qsec    vs    am  gear
##   <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 Merc 450SE   16.4     8.  276.  180.  3.07  4.07  17.4    0.    0.    3.
## 2 Merc 450SL   17.3     8.  276.  180.  3.07  3.73  17.6    0.    0.    3.
## 3 Merc 450SLC  15.2     8.  276.  180.  3.07  3.78  18.0    0.    0.    3.
## # ... with 1 more variable: carb <dbl>
```

## Chaining with %>%

when combining several function call together the command can be very hard to read

```
as.data.frame(select(filter(add_rownames(mtcars),mpg>=30),rowname,mpg,cyl,hp))
```

```
## Warning: Deprecated, use tibble::rownames_to_column() instead.
```

```
##       rowname  mpg cyl  hp
## 1      Fiat 128 32.4   4  66
## 2    Honda Civic 30.4   4  52
## 3 Toyota Corolla 33.9   4  65
## 4   Lotus Europa 30.4   4 113
```

You can chain commands together using the %>% operator.

$f(x) \%>\% g(y)$  is the same as  $g(f(x),y)$


```
add_rownames(mtcars) %>%
  filter(mpg>=30) %>%
  select(rowname,mpg,cyl,hp) %>%
  as.data.frame()
```

```
## Warning: Deprecated, use tibble::rownames_to_column() instead.
```

```
##       rowname  mpg cyl  hp
## 1      Fiat 128 32.4   4  66
## 2    Honda Civic 30.4   4  52
## 3 Toyota Corolla 33.9   4  65
## 4   Lotus Europa 30.4   4 113
```

The default is to put the left hand side as the first argument of the right hand side. You can use . as a placeholder to change this behaviour

```
filter(mtcars, cyl<8) %>%
  boxplot(mpg~cyl, data=.)
boxplot(mpg~cyl, data=mtcars, subset= cyl<8)
```



## Arrange

Use `arrange` to sort rows by value of a variable:

```
mtcars %>%
  arrange(mpg)
```

```
##      mpg  cyl  disp  hp drat   wt  qsec vs am gear carb
## 1  10.4    8  472.0  205  2.93 5.250 17.98  0  0    3    4
## 2  10.4    8  460.0  215  3.00 5.424 17.82  0  0    3    4
## 3  13.3    8  350.0  245  3.73 3.840 15.41  0  0    3    4
## 4  14.3    8  360.0  245  3.21 3.570 15.84  0  0    3    4
## 5  14.7    8  440.0  230  3.23 5.345 17.42  0  0    3    4
## 6  15.0    8  301.0  335  3.54 3.570 14.60  0  1    5    8
## 7  15.2    8  275.8  180  3.07 3.780 18.00  0  0    3    3
## 8  15.2    8  304.0  150  3.15 3.435 17.30  0  0    3    2
## 9  15.5    8  318.0  150  2.76 3.520 16.87  0  0    3    2
## 10 15.8    8  351.0  264  4.22 3.170 14.50  0  1    5    4
## 11 16.4    8  275.8  180  3.07 4.070 17.40  0  0    3    3
## 12 17.3    8  275.8  180  3.07 3.730 17.60  0  0    3    3
## 13 17.8    6  167.6  123  3.92 3.440 18.90  1  0    4    4
## 14 18.1    6  225.0  105  2.76 3.460 20.22  1  0    3    1
## 15 18.7    8  360.0  175  3.15 3.440 17.02  0  0    3    2
## 16 19.2    6  167.6  123  3.92 3.440 18.30  1  0    4    4
## 17 19.2    8  400.0  175  3.08 3.845 17.05  0  0    3    2
## 18 19.7    6  145.0  175  3.62 2.770 15.50  0  1    5    6
```



```
## 19 21.0 6 160.0 110 3.90 2.620 16.46 0 1 4 4
## 20 21.0 6 160.0 110 3.90 2.875 17.02 0 1 4 4
## 21 21.4 6 258.0 110 3.08 3.215 19.44 1 0 3 1
## 22 21.4 4 121.0 109 4.11 2.780 18.60 1 1 4 2
## 23 21.5 4 120.1 97 3.70 2.465 20.01 1 0 3 1
## 24 22.8 4 108.0 93 3.85 2.320 18.61 1 1 4 1
## 25 22.8 4 140.8 95 3.92 3.150 22.90 1 0 4 2
## 26 24.4 4 146.7 62 3.69 3.190 20.00 1 0 4 2
## 27 26.0 4 120.3 91 4.43 2.140 16.70 0 1 5 2
## 28 27.3 4 79.0 66 4.08 1.935 18.90 1 1 4 1
## 29 30.4 4 75.7 52 4.93 1.615 18.52 1 1 4 2
## 30 30.4 4 95.1 113 3.77 1.513 16.90 1 1 5 2
## 31 32.4 4 78.7 66 4.08 2.200 19.47 1 1 4 1
## 32 33.9 4 71.1 65 4.22 1.835 19.90 1 1 4 1
```

To break ties, just add more variables. Each additional variable will be used to break ties in the values of preceding ones:

```
arrange(mtcars,cyl,carb,gear) %>% select(cyl,carb,gear)
```

```
##      cyl carb gear
## 1      4     1    3
## 2      4     1    4
## 3      4     1    4
## 4      4     1    4
## 5      4     1    4
## 6      4     2    4
## 7      4     2    4
## 8      4     2    4
## 9      4     2    4
## 10     4     2    5
## 11     4     2    5
## 12     6     1    3
## 13     6     1    3
## 14     6     4    4
## 15     6     4    4
## 16     6     4    4
## 17     6     4    4
## 18     6     6    5
## 19     8     2    3
## 20     8     2    3
## 21     8     2    3
## 22     8     2    3
## 23     8     3    3
## 24     8     3    3
## 25     8     3    3
## 26     8     4    3
## 27     8     4    3
## 28     8     4    3
## 29     8     4    3
## 30     8     4    3
## 31     8     4    5
## 32     8     8    5
```

Use `desc()` to sort in descending order:

```
mtcars %>%
  arrange(desc(mpg))
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## 1	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## 2	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## 3	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## 4	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## 5	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
## 6	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
## 7	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## 8	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
## 9	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
## 10	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
## 11	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## 12	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2
## 13	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## 14	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## 15	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
## 16	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
## 17	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
## 18	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## 19	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## 20	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
## 21	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## 22	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## 23	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
## 24	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## 25	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## 26	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## 27	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
## 28	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## 29	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## 30	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## 31	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## 32	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4

## Mutate

To add new variables based on some calculation, possibly using the value of other variables, use `mutate`:

```
mutate(mtcars, displ_1 = disp / 61.0237)
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	displ_1
## 1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4	2.621932
## 2	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4	2.621932
## 3	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1	1.769804
## 4	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1	4.227866
## 5	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2	5.899347
## 6	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1	3.687092
## 7	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4	5.899347
## 8	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2	2.403984
## 9	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2	2.307300
## 10	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4	2.746474

```
## 11 17.8 6 167.6 123 3.92 3.440 18.90 1 0 4 4 2.746474
## 12 16.4 8 275.8 180 3.07 4.070 17.40 0 0 3 3 4.519556
## 13 17.3 8 275.8 180 3.07 3.730 17.60 0 0 3 3 4.519556
## 14 15.2 8 275.8 180 3.07 3.780 18.00 0 0 3 3 4.519556
## 15 10.4 8 472.0 205 2.93 5.250 17.98 0 0 3 4 7.734700
## 16 10.4 8 460.0 215 3.00 5.424 17.82 0 0 3 4 7.538055
## 17 14.7 8 440.0 230 3.23 5.345 17.42 0 0 3 4 7.210313
## 18 32.4 4 78.7 66 4.08 2.200 19.47 1 1 4 1 1.289663
## 19 30.4 4 75.7 52 4.93 1.615 18.52 1 1 4 2 1.240502
## 20 33.9 4 71.1 65 4.22 1.835 19.90 1 1 4 1 1.165121
## 21 21.5 4 120.1 97 3.70 2.465 20.01 1 0 3 1 1.968088
## 22 15.5 8 318.0 150 2.76 3.520 16.87 0 0 3 2 5.211090
## 23 15.2 8 304.0 150 3.15 3.435 17.30 0 0 3 2 4.981671
## 24 13.3 8 350.0 245 3.73 3.840 15.41 0 0 3 4 5.735477
## 25 19.2 8 400.0 175 3.08 3.845 17.05 0 0 3 2 6.554830
## 26 27.3 4 79.0 66 4.08 1.935 18.90 1 1 4 1 1.294579
## 27 26.0 4 120.3 91 4.43 2.140 16.70 0 1 5 2 1.971365
## 28 30.4 4 95.1 113 3.77 1.513 16.90 1 1 5 2 1.558411
## 29 15.8 8 351.0 264 4.22 3.170 14.50 0 1 5 4 5.751864
## 30 19.7 6 145.0 175 3.62 2.770 15.50 0 1 5 6 2.376126
## 31 15.0 8 301.0 335 3.54 3.570 14.60 0 1 5 8 4.932510
## 32 21.4 4 121.0 109 4.11 2.780 18.60 1 1 4 2 1.982836
```

You can add as many variables as you want, and even base their value on any preceding column:

```
mutate(mtcars, displ_l = disp / 61.0237, wt_kg = wt / 2.2, wt_rt = hp / wt_kg)
```

```
##      mpg cyl  disp  hp drat   wt  qsec vs am gear carb  displ_l    wt_kg
## 1  21.0   6 160.0 110 3.90 2.620 16.46 0  1   4    4 2.621932 1.1909091
## 2  21.0   6 160.0 110 3.90 2.875 17.02 0  1   4    4 2.621932 1.3068182
## 3  22.8   4 108.0  93 3.85 2.320 18.61 1  1   4    1 1.769804 1.0545455
## 4  21.4   6 258.0 110 3.08 3.215 19.44 1  0   3    1 4.227866 1.4613636
## 5  18.7   8 360.0 175 3.15 3.440 17.02 0  0   3    2 5.899347 1.5636364
## 6  18.1   6 225.0 105 2.76 3.460 20.22 1  0   3    1 3.687092 1.5727273
## 7  14.3   8 360.0 245 3.21 3.570 15.84 0  0   3    4 5.899347 1.6227273
## 8  24.4   4 146.7  62 3.69 3.190 20.00 1  0   4    2 2.403984 1.4500000
## 9  22.8   4 140.8  95 3.92 3.150 22.90 1  0   4    2 2.307300 1.4318182
## 10 19.2   6 167.6 123 3.92 3.440 18.30 1  0   4    4 2.746474 1.5636364
## 11 17.8   6 167.6 123 3.92 3.440 18.90 1  0   4    4 2.746474 1.5636364
## 12 16.4   8 275.8 180 3.07 4.070 17.40 0  0   3    3 4.519556 1.8500000
## 13 17.3   8 275.8 180 3.07 3.730 17.60 0  0   3    3 4.519556 1.6954545
## 14 15.2   8 275.8 180 3.07 3.780 18.00 0  0   3    3 4.519556 1.7181818
## 15 10.4   8 472.0 205 2.93 5.250 17.98 0  0   3    4 7.734700 2.3863636
## 16 10.4   8 460.0 215 3.00 5.424 17.82 0  0   3    4 7.538055 2.4654545
## 17 14.7   8 440.0 230 3.23 5.345 17.42 0  0   3    4 7.210313 2.4295455
## 18 32.4   4 78.7  66 4.08 2.200 19.47 1  1   4    1 1.289663 1.0000000
## 19 30.4   4 75.7  52 4.93 1.615 18.52 1  1   4    2 1.240502 0.7340909
## 20 33.9   4 71.1  65 4.22 1.835 19.90 1  1   4    1 1.165121 0.8340909
## 21 21.5   4 120.1  97 3.70 2.465 20.01 1  0   3    1 1.968088 1.1204545
## 22 15.5   8 318.0 150 2.76 3.520 16.87 0  0   3    2 5.211090 1.6000000
## 23 15.2   8 304.0 150 3.15 3.435 17.30 0  0   3    2 4.981671 1.5613636
## 24 13.3   8 350.0 245 3.73 3.840 15.41 0  0   3    4 5.735477 1.7454545
## 25 19.2   8 400.0 175 3.08 3.845 17.05 0  0   3    2 6.554830 1.7477273
## 26 27.3   4 79.0  66 4.08 1.935 18.90 1  1   4    1 1.294579 0.8795455
## 27 26.0   4 120.3  91 4.43 2.140 16.70 0  1   5    2 1.971365 0.9727273
```

```
## 28 30.4 4 95.1 113 3.77 1.513 16.90 1 1 5 2 1.558411 0.6877273
## 29 15.8 8 351.0 264 4.22 3.170 14.50 0 1 5 4 5.751864 1.4409091
## 30 19.7 6 145.0 175 3.62 2.770 15.50 0 1 5 6 2.376126 1.2590909
## 31 15.0 8 301.0 335 3.54 3.570 14.60 0 1 5 8 4.932510 1.6227273
## 32 21.4 4 121.0 109 4.11 2.780 18.60 1 1 4 2 1.982836 1.2636364
##      wt_rt
## 1  92.36641
## 2  84.17391
## 3  88.18966
## 4  75.27216
## 5 111.91860
## 6  66.76301
## 7 150.98039
## 8  42.75862
## 9  66.34921
## 10 78.66279
## 11 78.66279
## 12 97.29730
## 13 106.16622
## 14 104.76190
## 15 85.90476
## 16 87.20501
## 17 94.66791
## 18 66.00000
## 19 70.83591
## 20 77.92916
## 21 86.57201
## 22 93.75000
## 23 96.06987
## 24 140.36458
## 25 100.13004
## 26 75.03876
## 27 93.55140
## 28 164.30932
## 29 183.21767
## 30 138.98917
## 31 206.44258
## 32 86.25899
```

To only keep the newly created variables, use `transmute`:

```
transmute(mtcars, displ_l = disp / 61.0237)
```

```
##      displ_l
## 1  2.621932
## 2  2.621932
## 3  1.769804
## 4  4.227866
## 5  5.899347
## 6  3.687092
## 7  5.899347
## 8  2.403984
## 9  2.307300
## 10 2.746474
## 11 2.746474
```

```
## 12 4.519556
## 13 4.519556
## 14 4.519556
## 15 7.734700
## 16 7.538055
## 17 7.210313
## 18 1.289663
## 19 1.240502
## 20 1.165121
## 21 1.968088
## 22 5.211090
## 23 4.981671
## 24 5.735477
## 25 6.554830
## 26 1.294579
## 27 1.971365
## 28 1.558411
## 29 5.751864
## 30 2.376126
## 31 4.932510
## 32 1.982836
```

## Summarise

`summarise` calculates a single value using a set of variable values:

```
summarise(mtcars, mean(mpg))
```

```
##      mean(mpg)
## 1  20.09062
```

This is the same as, and more wordy than, just using the same function on a data frame column:

```
mean(mtcars$mpg)
```

```
## [1] 20.09062
```

```
summarise(mtcars, mpg=mean(mpg))
```

```
##      mpg
## 1 20.09062
```

Its benefits get more obvious when calculating multiple summaries or when the calculation is based on more than one column:

```
summarise(mtcars, mpg=mean(mpg), wt=median(wt))
```

```
##      mpg      wt
## 1 20.09062 3.325
```

```
summarise(mtcars, mpg=mean(mpg), wt=median(wt), rat=mpg/wt)
```

```
##      mpg      wt      rat
## 1 20.09062 3.325 6.042293
```

You can use any function that takes in a vector of values and returns a scalar (i.e., “aggregates”): `mean`, `median`, `max`, etc.

But where `summarise` really comes handy is when we want to calculate it for groups of observations. This is done by first applying the `group_by` verb and then feed it into `summarise`. For instance, to calculate the mean gas mileage for each engine size:

```
mtcars %>%  
  group_by(cyl) %>%  
  summarise(mpg = mean(mpg), wt=median(wt))
```

```
## # A tibble: 3 x 3  
##   cyl  mpg   wt  
##   <dbl> <dbl> <dbl>  
## 1     4.  26.7  2.20  
## 2     6.  19.7  3.22  
## 3     8.  15.1  3.76
```

Note that the grouping variable are included in the result. Using more than one grouping variable will split by each one in turn:

```
mtcars %>%  
  group_by(am, cyl) %>%  
  summarise(mpg = mean(mpg), wt=median(wt))
```

```
## # A tibble: 6 x 4  
## # Groups:   am [?]  
##   am  cyl  mpg   wt  
##   <dbl> <dbl> <dbl> <dbl>  
## 1  0.    4.  22.9  3.15  
## 2  0.    6.  19.1  3.44  
## 3  0.    8.  15.0  3.81  
## 4  1.    4.  28.1  2.04  
## 5  1.    6.  20.6  2.77  
## 6  1.    8.  15.4  3.37
```

Each summarizing “rolls up” one grouping, starting from the bottom:

```
cars_am_cyl <- mtcars %>%  
  group_by(am, cyl) %>%  
  summarise(mpg = mean(mpg), wt=median(wt))  
cars_am_cyl
```

```
## # A tibble: 6 x 4  
## # Groups:   am [?]  
##   am  cyl  mpg   wt  
##   <dbl> <dbl> <dbl> <dbl>  
## 1  0.    4.  22.9  3.15  
## 2  0.    6.  19.1  3.44  
## 3  0.    8.  15.0  3.81  
## 4  1.    4.  28.1  2.04  
## 5  1.    6.  20.6  2.77  
## 6  1.    8.  15.4  3.37
```

```
cars_am_cyl %>%  
  summarise(mpg = mean(mpg), wt=median(wt))
```

```
## # A tibble: 2 x 3  
##   am  mpg   wt  
##   <dbl> <dbl> <dbl>  
## 1  0.  19.0  3.44
```

## 2 1. 21.3 2.77